


Improving System Performance

Computer Engineering 1

Agenda

- Motivation
- Aspects of Optimization
- Bus Architectures
- Instruction Set Architectures (ISA)
- Pipelining
- Parallel Computing



■ Intel Pentium E5800

- 3.2 GHz
- 2 Cores
- 2 MB cache
- Released Q4 2010
- **PassMark: 1184**
- **TDP: 65 W**

SPEED!




■ Intel Celeron G4930

- 3.2 GHz
- 2 Cores
- 2 MB cache
- Q2 2019
- **PassMark: 2628**
- **TDP: 54 W**

■ AMD Ryzen R1102G


- 1.2 GHz
- 2 Cores
- 4 MB cache
- Q1 2020
- **PassMark: ~1600**
- **TDP: 6 W**

LOW POWER!



Motivation

■ Growth in processor performance




Source: John L. Hennessy, David A. Patterson (2017). Computer Architecture: A Quantitative Approach. 6th edition.

■ **At the end of this lesson you will be able**

- to explain different types of bus architectures
- to understand the difference between von Neumann and Harvard architecture
- to understand RISC and CISC paradigms
- to describe the idea of pipelining
- to calculate processing performance improvement through pipelining
- to describe the basics of parallel computing

■ ‘Wishlist of Optimization’

Optimizing for:



Drawbacks on:

Power, cost, chip area

Speed, reliability


Speed, cost

Chip area, cost, speed

Power, cost, lifetime

- This lecture: Increasing system/CPU speed
- Some other topics: Covered in ‘Microcomputer Systems 1’ (MC1)

How to Increase System Speed?



CPU improvements:

- Clock Speed
- Cache Memory
- Multiple Cores / Threads
- Pipelined Execution
- Branch Prediction
- Out of Order Execution
- Instruction Set Architecture

External factors:

- Better Compilers
- Better Algorithms

System level factors:

- Special Purpose Units
e.g. Crypto, Video, AI
- Multiple Processors
- Bus Architecture
e.g. von Neumann / Harvard
- Faster components (e.g. SSDs, 1000Base-T, etc.)


How is the connection between the CPU and memory organized?

A single system bus: **simple, cheaper** **slower**

Multiple system buses: **faster** **more complexity**


■ von Neumann Architecture

- Same memory holds program and data
- Single bus system between CPU and memory




■ Harvard Architecture

- “Mark I” at Harvard University (Howard Aiken, 1939-44)
- Separate memories for program and data
- Two sets of address/data buses between CPU and memory




■ von Neumann Architecture

- Example: ARM Cortex-M0



■ Harvard Architecture

- Example: ARM Cortex-M3/M4



Powerful, complex instructions?

or

Simple, highly optimizable instructions?

(or something in between?)

■ First CPUs

- Small address bus, small data bus (8 bit)
- Few registers, no cache, pipelining, co-processors, multi-core
- Machine code as finite state machine: Shift, multiply, etc. in loops

■ Consequences for machine code

- Differing complexity
- Different length
- Variable execution time
- Individual addressing / arguments

} for different instructions

→ **More and more complex instructions (CISC)**

- **John Cocke (IBM, 1974) proved that**
 - 80% of work was done using only 20% of the instructions
 - Most (complex) instructions were not used at all by compilers
- **RISC (Reduced Instruction Set Computer)**
 - Few instructions, unique instruction format
 - Fast decoding, simple addressing
 - less hardware → allows higher clock rates
 - more chip space for registers (up to 256!)
 - Load-store architecture reduces memory accesses, CPU works at full-speed on registers



Acorn Archimedes A3000
Source: Wikipedia Commons

Remark: Word “CISC” was introduced with “RISC”

Example: **Balance = Balance + Credit**

■ RISC

- Load / Store Architecture
- Data processing instructions only available on registers


```
LDR    R0,=Credit
LDR    R1,[R0]
LDR    R0,=Balance
LDR    R3,[R0]
ADDS   R2,R1,R3
STR    R2,[R0]
```

■ CISC

- One of the operands of an instruction may directly be a memory location

```
MOV    AX, [Credit]
ADD    [Balance], AX
```

■ Instruction / opcode complexity



■ RISC advantages

- Simple and fast instruction decoding
- More registers/cache on silicon area (less memory access)
- Several data paths possible (superscalar computers)
- Higher clock frequencies
- Effective compiler optimization with limited, generic instructions
- Easy pipelining, shorter pipelines (instruction size / duration)


■ CISC advantages

- Less program memory needed with complex instructions
- Short programs may work faster with less memory accesses

■ CISC and RISC more and more indistinguishable


- Modern x86 CPUs convert instructions (CISC to RISC)

■ ARM Cortex-M3 sequence:



Source: ARM


■ Idea: Why not fetch the next instruction, while the current one decodes?



■ Timings and definitions


- Example: Cortex-M3

- | | | |
|----------------------|---|-------------------|
| - fe: fetch | Read instruction | 3 ns ¹ |
| - de: decode | Decode instruction, read register or memory | 4 ns ¹ |
| - ex: execute | Execute instruction, write back result | 5 ns ¹ |




¹timings estimated

■ Sequential execution



■ Pipelined execution



■ Instructions per second

- Without pipelining

$$\text{Instructions per second} = \frac{1}{\text{Instruction delay}}$$

- With pipelining
 - Pipeline needs to be filled first
 - After filling, instructions are executed after every stage

$$\text{Instructions per second} = \frac{1}{\text{Max stage delay}}$$

■ Example Cortex-M3

- Without pipelining: $1/12\text{ns} = 83,3$ million instructions per second
- With pipelining: $1/5\text{ns} = 200$ million instructions per second

■ Advantages

- All stages are set to the same execution time
- Massive performance gain
- Simpler hardware at each stage allows for a higher clock rate

■ Disadvantages

- A blocking stage blocks whole pipeline
- Multiple stages may need to have access to the memory at the same time

■ General

- Number of execution stages is design decision
- Typically 3 (ARM Cortex M4) – 20 stages (Pentium 4)


■ Optimal Pipelining

- All operations here are on registers (single cycle execution)
- In this example it takes 6 clock cycles to execute 6 instructions
- Clock cycles per instruction (CPI) = 1

Cycle		1	2	3	4	5	6	7	8	9
Operation										
<i>ADD</i>	fe	de	ex							
<i>SUB</i>		fe	de	ex						
<i>ORR</i>			fe	de	ex					
<i>AND</i>				fe	de	ex				
<i>ORR</i>					fe	de	ex			
<i>EOR</i>						fe	de	ex		

■ Special situation: LDR


- In this example it takes 7 clock cycles to execute 6 instructions
- Read cycle must complete on the bus before LDR instruction can complete
- Next 2 instructions must wait one pipeline cycle (S = stall)
- Clock cycles per instruction (CPI) = 1.2



■ Control Hazards

- Branch / jump decisions occur in stage 3 (ex)
- Worst case scenario – conditional branch taken:

00000000	4288	CMP	R0,R1	; R0 > R1 ?
00000002	D802	BHI	go_on	; if higher -> go_on
00000004	000A	MOVS	R2,R1	; otherwise exchange regs
00000006	0001	MOVS	R1,R0	
00000008	0010	MOVS	R0,R2	
0000000A	2305	go_on	MOVS	R3,#5
0000000C	...	...		



■ Reduce control hazards

- **Loop fusion** reduces control hazards
- Simple example for illustration

```
/* Two loops → double number of control hazards*/  
for (i = 0; i < N; i++) {  
    a[i] = 1/b[i] * c[i];  
}  
for (i = 0; i < N; i++) {  
    d[i] = a[i] + c[i];  
}  
  
/* Better performance: Fusion into one loop */  
for (i = 0; i < N; i++) {  
    a[i] = 1/b[i] * c[i];  
    d[i] = a[i] + c[i];  
}
```

2*N Conditional
Branches

N Conditional
Branches

■ Ideas to further improve pipelining:

- Branch prediction:
 - Store last decision made for each conditional branch
→ probability is high that the same decision is taken again
- Instruction prefetch:
 - Fetch several instructions in advance
→ better use of the system bus
→ possibility of 'Out of Order Execution'
- Out of Order Execution:
 - If one instruction stalls (e.g. a LDR from slow memory), it might be possible to already execute the next instruction

■ Limits of optimization

- Complex Optimizations → severe security problems (e.g. Out of Order Execution)
- Instructions (speculatively) executed, that would throw access violations under 'In Order' circumstances.
- 'Meltdown' and 'Spectre' attacks: allow a process to access the data of another process.



Streaming/Vector Processing

- One instruction processes multiple data items simultaneously

Multithreading

- Multiple programs/threads share a single CPU

Multicore Processors

- One processor contains multiple CPU cores

Multiprocessor Systems

- A computer system contains multiple processors

■ Instructions working on more than one data item

		Data streams	
		Single	Multiple
Instruction streams	Single	SISD : Single processor with one instruction, data stream	SIMD : Data Vectors, all data streams react to one instruction
	Multiple	(no examples)	MIMD : Multiprocessor with SIMD-instructions


Source: David A. Patterson, John L. Hennessy

■ SIMD examples (x86):

MMX (Multimedia Extension), AMD 3D-Now,
SSE (Streaming SIMD Extensions),
AVX (Advanced Vector Extensions)

■ Multithreading

- Scheduler: Assigns time slots to programs/threads
- Programs/threads only **seem** to run in parallel (1 CPU)




Remark: More on Multithreading in MC1

■ Parallelism on CPU / processor level:


- Multicore processor

- All on one chip
- Less traffic (cores integrated on one chip)
- Possibility to share memories on-chip
- Cheaper




- Multiprocessor

- Multiple Chips
- Longer distances between CPUs
- More expensive



Source: <https://www.queryhome.com/tech/106320/what-the-difference-between-multicore-and-multiprocessor>

- **CPU clock frequency not the only factor**
- **Different ways to increase processor performance**
 - RISC / CISC
 - Pipelining
 - Out-of-order execution
 - Parallel computing
- **Today's CPUs combine most of these technologies**
- **Other components with influence**
 - SSD, memory, interfaces, algorithms ...



Source: <https://researchcomputingservices.github.io/parallel-computing/02-speedup-limitations/>